

Concept Evaluator — What We Changed and Why

Prepared for Jon and Mike — 2026-06-15

This is a plain-English summary of the upgrades made to the business-concept evaluation system. The goal of the system is unchanged: ****find 2–3 need-to-have businesses we can launch on low cash and grow fast, scored the way real expert advisors would score them.**** The changes below fix the things that were getting in the way of that.

The core problem we fixed

The old system was quietly built to **hold scores down**. Three things combined to do it:

1. The expert reviewers were told to back any high score with hard evidence (named competitors, real numbers) — but they had **no way to look anything up** (no web access). So they couldn't justify a high score honestly, and defaulted to mediocre.
2. Anything not backed by a hard number was **capped at 6 out of 10**.
3. The bar to "pass" was an average above 8.0 with no score below 7.

Put together, the math was nearly impossible: the system's own notes admit every idea clustered at 6.6–7.2 and **nothing ever passed**. The only way to "pass" was to fudge numbers — which the rules then forbade. So the system was stuck, and the scores felt artificially suppressed. They were.

The fix: give the market-facing reviewers real web research, make the evidence standard match the fact that we're at the concept stage, and make scoring honest in **both** directions — a genuinely strong idea should earn an 8 or 9, and we say so.

The five upgrades

1. We now score at the right stage (concept stage)

A new ground-rule doc states plainly: we are evaluating **concepts**, before customer

interviews and before anything is built. Reviewers can no longer dock points because we "haven't talked to customers yet" or "haven't built it." Those become the ****list of cheap things to test next****, not penalties. A strong idea backed by a named comparable and sound math can now score an 8 or 9 — which is how a real advisor would score it.

2. Competition gets a real, dedicated analyst

Competition was one paragraph inside the marketing review, written by someone who couldn't browse the web. We added a **dedicated Competitive Analyst with web access** whose whole job is the competitive teardown: name the real players, pull their pricing and funding, judge whether they're genuinely sleepy and beatable or actually dangerous, check whether a big platform could move into the niche, and war-game what the incumbents do when we start winning. Competition is now one of the two heaviest-weighted factors — because out-positioning sleepy incumbents is our edge.

3. "Must-have" no longer means "regulated"

We were over-indexing on compliance. Now a must-have is defined by its ****forcing function**** — what breaks, and how fast, if the customer doesn't buy — and there are **five equally valid kinds**:

1. A regulation or compliance mandate forces the buy.
2. A customer, prime contractor, insurer, lender, or franchisor requires it.
3. It's a **critical input to their production** — without it, their line stops or their product can't ship. (Your raw-material example — now treated as exactly as strong as a compliance mandate.)
4. Without it they lose money, customers, or the ability to operate.
5. It's effectively mandatory risk mitigation that responsible operators won't skip.

Regulation is one favorite path, never required. A raw material a factory can't run without scores just as high.

4. Regulation is no longer an over-weighted drag

The old setup had a perverse bug: the Legal reviewer scored "more regulation = worse

score," which **punished exactly the forced-demand businesses we like**. Now Legal runs a **risk gate** instead — it flags exposure as None / Minor / Serious-but-Manageable / Fatal. A regulated business with a clear, navigable compliance path is **not** marked down for being regulated; it just carries a due-diligence checklist. Only a genuinely fatal legal problem can sink a concept. The **upside** of regulation (it forces demand) is now credited where it belongs — in the must-have score.

5. Plain-English artifacts you can actually share

Everything used to live as dense markdown in GitHub, which is hard for you to read and hard to share with Mike. Now every concept we surface comes with:

- A **Concept Dossier** in plain English (what it is, who pays and why they can't say no, the money in simple terms, why we win, the competition, the top risks, what to test next, and the scores explained in plain sentences).
- A **one-page Scorecard**.
- A **master comparison sheet** ranking the whole field at a glance.

These are generated as **Word (.docx)**, **PDF**, and **text (.txt)** files, saved in the project, **pushed to your Google Drive**, and sent to you as attachments — no GitHub required. (This very document was produced by that new pipeline as a test.)

How scoring works now (and what "passing" means)

We dropped the rigid "must average above 8.0" gate. Instead each concept gets an **Objectives Score out of 100** and an honest **tier**:

- **A — Advisor-Ready**: strong enough to put in front of expert advisors and defend.
- **B — Promising**: worth a deeper look before committing.
- **C — Pass**: drop it.

The deliverable is always the **honest top 2–3 of the field, ranked**. If the best we find are B-tier, we'll say so and still hand you the top 2–3 with exactly what would move them up — because you're going to compare them against real experts anyway. We do not inflate to manufacture a winner, and we no longer suppress a real winner out

of caution.

What to do next

When you're ready to run a search, start a fresh session and use one of the run modes (the simplest is the `/goal` command in doc 09). It will brainstorm widely across all five must-have types, run the six-member council on the survivors, and hand you plain-English dossiers for the top 2–3 in your Google Drive. Then take those to your real advisors and compare — which was the point all along.